

University of Tehran  
Department of ECE  
Neural Networks & Deep Learning  
MP1

Members: Milad Mohammadi & Tohid Abdi  
IDs: 810100462 & 810100410  
Date: April 24, 2022

## Questions List Table

<b>1</b>	<b>CNN Classification</b>	<b>3</b>
1.1	Convolutional Layers . . . . .	5
1.2	Batch Normalization and Pooling . . . . .	9
1.3	Dropout . . . . .	14
1.4	Early Stop . . . . .	19
<b>2</b>	<b>Transfer Learning</b>	<b>21</b>
2.1	Model Properties . . . . .	21
2.2	Transfer Learning . . . . .	21
2.3	Implementation . . . . .	24
<b>3</b>	<b>Segmentation</b>	<b>27</b>
<b>4</b>	<b>Object Detection</b>	<b>36</b>
4.1	YOLO v1 and v2 . . . . .	36
4.2	YOLO 4 & YOLO 5: . . . . .	36
4.3	Two-stage vs One-stage Detectors: . . . . .	37
4.4	YOLO v5 implementation: . . . . .	37

# 1 CNN Classification

import essential libraries:

```
[1]: from keras.datasets import cifar10
from sklearn.model_selection import train_test_split
import random
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from keras.layers import BatchNormalization
import pandas as pd
from sklearn.metrics import confusion_matrix, precision_recall_fscore_support
import keras.utils.vis_utils
from importlib import reload
import pydot
reload(keras.utils.vis_utils)
keras.utils.vis_utils.pydot = pydot
from tensorflow.keras.utils import to_categorical
```

import cifar10:

```
[2]: (X_train_full, y_train_full), (X_test, y_test) = cifar10.load_data()
```

normalize data:

```
[3]: X_train_normal = X_train_full/255
X_test_normal = X_test/255
```

train-test split:

```
[4]: X_train, X_val, y_train, y_val = train_test_split(X_train_normal, y_train_full,
↳test_size=0.2, random_state = 20)
```

transform dependent variables to one hot:

```
[5]: y_train_onehot = to_categorical(y_train, 10)
y_test_onehot = to_categorical(y_test, 10)
y_val_onehot = to_categorical(y_val, 10)
```

set random seed:

```
[6]: np.random.seed(0)
tf.random.set_seed(0)
```

make call back for early stopping:

```
[7]: callback = tf.keras.callbacks.EarlyStopping(monitor='val_accuracy', patience=4,
→restore_best_weights=True)
```

define functions to compiling data, plotting accuracies and losses, evaluating test data and drawing model architecture:

```
[8]: def model_compile(loss, optimizer, batch_s):

    model.compile(loss = loss,
                  optimizer = optimizer,
                  metrics = ['accuracy'])

    global model_history
    model_history = model.fit(X_train, y_train_onehot, epochs=500, batch_size =
→batch_s, validation_data=(X_val, y_val_onehot), callbacks = [callback])
```

```
[9]: def plot(x):

    loss = model_history.history.copy()
    loss.pop('accuracy')
    loss.pop('val_accuracy')
    acc = model_history.history.copy()
    acc.pop('loss')
    acc.pop('val_loss')

    pd.DataFrame(loss).plot(figsize=(8,5))
    plt.grid(True)
    plt.show()

    pd.DataFrame(acc).plot(figsize=(8,5))
    plt.grid(True)
    plt.ylim(0,1)
    plt.show()
```

```
[10]: def test_evaluate(x):

    test_loss = x.evaluate(X_test_normal, y_test_onehot)[0]
    test_accuracy = x.evaluate(X_test_normal, y_test_onehot)[1]
    y_prob = x.predict(X_test_normal)
    y_pred = y_prob.argmax(axis=-1)
    cnf = confusion_matrix(y_test, y_pred)
    prf = precision_recall_fscore_support(y_test, y_pred, average = 'macro')

    print('test loss: ' + str(test_loss) + '\n\ntest accuracy: ' +
→str(test_accuracy) + '\n\nconfusion matrix: \n' + str(cnf) + '\n\nf1-score: '
→+ str(prf[2]) + '\n\nrecall: ' + str(prf[1]) + '\n\nprecision: ' + str(prf[0]))
```

```
[11]: def model_architecture(x):
        return tf.keras.utils.plot_model(
            x,
            show_shapes=True,
            show_dtype=True,
            show_layer_names=False,
            rankdir="LR",
            expand_nested=True,
            dpi=96,
            layer_range=None,
            show_layer_activations=True,
        )
```

## 1.1 Convolutional Layers

model with solely convolution layers:

```
[12]: model = keras.models.Sequential()
model.add(layers.Conv2D(filters=16, kernel_size=(3,3), strides=1,
    ↳activation='relu', padding='same', input_shape=(32, 32, 3)))
model.add(layers.Conv2D(filters=16, kernel_size=(3,3), strides=1,
    ↳kernel_initializer='he_uniform', activation='relu', padding='same'))

model.add(layers.Conv2D(filters=32, kernel_size=(3,3), strides=1,
    ↳activation='relu', padding='same'))
model.add(layers.Conv2D(filters=32, kernel_size=(3,3), strides=1,
    ↳activation='relu', padding='same'))

model.add(layers.Conv2D(filters=64, kernel_size=(3,3), strides=1,
    ↳activation='relu', padding='same'))
model.add(layers.Conv2D(filters=64, kernel_size=(3,3), strides=1,
    ↳activation='relu', padding='same'))

model.add(keras.layers.Flatten())
model.add(keras.layers.Dense(64, activation = 'relu'))
model.add(keras.layers.Dense(10, activation = 'softmax'))

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 32, 32, 16)	448
conv2d_1 (Conv2D)	(None, 32, 32, 16)	2320

conv2d_2 (Conv2D)	(None, 32, 32, 32)	4640
conv2d_3 (Conv2D)	(None, 32, 32, 32)	9248
conv2d_4 (Conv2D)	(None, 32, 32, 64)	18496
conv2d_5 (Conv2D)	(None, 32, 32, 64)	36928
flatten (Flatten)	(None, 65536)	0
dense (Dense)	(None, 64)	4194368
dense_1 (Dense)	(None, 10)	650

```
=====
Total params: 4,267,098
Trainable params: 4,267,098
Non-trainable params: 0
-----
```

```
[13]: model_compile('categorical_crossentropy', 'adam', 32)
      plot(model_history.history)
      test_evaluate(model)
      model_architecture(model)
```

```
Epoch 1/500
1250/1250 [=====] - 252s 201ms/step - loss: 1.5978 -
accuracy: 0.4211 - val_loss: 1.2737 - val_accuracy: 0.5414
...
Epoch 9/500
1250/1250 [=====] - 255s 204ms/step - loss: 0.1287 -
accuracy: 0.9543 - val_loss: 2.1546 - val_accuracy: 0.6258
```




Figure 1: Loss and validation loss for model with convolutional layers




Figure 2: Accuracy and validation accuracy for model with convolutional layers

```
313/313 [=====] - 11s 36ms/step - loss: 1.1573 -
accuracy: 0.6429
313/313 [=====] - 11s 36ms/step - loss: 1.1573 -
accuracy: 0.6429
test loss: 1.1573235988616943
```

```
test accuracy: 0.6428999900817871
```

```
confusion matrix:
[[728  24  47  26  21   6   3  13  98  34]
 [ 41 800   4  13   6   3   3   7  36  87]
 [109  11 467  82 128  63  41  61  29   9]
 [ 41  31  60 464  80 146  50  73  33  22]
 [ 34   3  88  78 594  42  33 100  23   5]
 [ 25  11  45 182  57 535  23  98  19   5]
 [ 19  18  79  99  68  28 635  22  21  11]
 [ 19   3  29  36  73  56   1 753  13  17]
 [127  50   7  13  12   5   1  13 730  42]
 [ 64 109   5  20   7   6   0  23  43 723]]
```

```
f1-score: 0.6413638859917244
```



recall: 0.6428999999999999

precision: 0.6453779023882527

[13]:




Figure 3: Architecture of model with convolutional layers

## 1.2 Batch Normalization and Pooling

**batch normalization:** Batch-Normalization (BN) is an algorithmic method which makes the training of Deep Neural Networks (DNN) faster and more stable. It consists of normalizing activation vectors from hidden layers using the first and the second statistical moments (mean and variance) of the current batch. This normalization step is applied right before (or right after) the nonlinear function.

**pooling:** Pooling layers provide an approach to down sampling feature maps by summarizing the presence of features in patches of the feature map. Two common pooling methods are average pooling and max pooling that summarize the average presence of a feature and the most activated presence of a feature respectively.

implementing model with batch normalization and pooling:

```
[14]: model = keras.models.Sequential()
model.add(layers.Conv2D(filters=16, kernel_size=(3,3), strides=1,
    ↪activation='relu', padding='same', input_shape=(32, 32, 3)))
model.add(layers.Conv2D(filters=16, kernel_size=(3,3), strides=1,
    ↪activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(layers.MaxPooling2D(2, 2))

model.add(layers.Conv2D(filters=32, kernel_size=(3,3), strides=1,
    ↪activation='relu', padding='same'))
model.add(layers.Conv2D(filters=32, kernel_size=(3,3), strides=1,
    ↪activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(layers.MaxPooling2D(2, 2))

model.add(layers.Conv2D(filters=64, kernel_size=(3,3), strides=1,
    ↪activation='relu', padding='same'))
```

```

model.add(layers.Conv2D(filters=64, kernel_size=(3,3), strides=1,
    ↪activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(layers.MaxPooling2D(2, 2))

model.add(keras.layers.Flatten())
model.add(keras.layers.Dense(64, activation = 'relu'))
model.add(BatchNormalization())
model.add(keras.layers.Dense(10, activation = 'softmax'))

model.summary()

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 32, 32, 16)	448
conv2d_7 (Conv2D)	(None, 32, 32, 16)	2320
batch_normalization (Batch Normalization)	(None, 32, 32, 16)	64
max_pooling2d (MaxPooling2D)	(None, 16, 16, 16)	0
conv2d_8 (Conv2D)	(None, 16, 16, 32)	4640
conv2d_9 (Conv2D)	(None, 16, 16, 32)	9248
batch_normalization_1 (Batch Normalization)	(None, 16, 16, 32)	128
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 32)	0
conv2d_10 (Conv2D)	(None, 8, 8, 64)	18496
conv2d_11 (Conv2D)	(None, 8, 8, 64)	36928
batch_normalization_2 (Batch Normalization)	(None, 8, 8, 64)	256
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 64)	0
flatten_1 (Flatten)	(None, 1024)	0

dense_2 (Dense)	(None, 64)	65600
batch_normalization_3 (Batch Normalization)	(None, 64)	256
dense_3 (Dense)	(None, 10)	650

```
=====
Total params: 139,034
Trainable params: 138,682
Non-trainable params: 352
-----
```

```
[15]: model_compile('categorical_crossentropy', 'adam', 32)
      plot(model_history.history)
      test_evaluate(model)
      model_architecture(model)
```

```
Epoch 1/500
1250/1250 [=====] - 74s 58ms/step - loss: 1.3735 -
accuracy: 0.5133 - val_loss: 1.2330 - val_accuracy: 0.5631
...
Epoch 17/500
1250/1250 [=====] - 68s 54ms/step - loss: 0.1683 -
accuracy: 0.9406 - val_loss: 1.0347 - val_accuracy: 0.7447
```




Figure 4: Loss and validation loss for CNN model with batch normalization and Max pooling





Figure 5: Accuracy and validation accuracy for CNN model with batch normalization and Max pooling

```
313/313 [=====] - 4s 13ms/step - loss: 0.9102 -
accuracy: 0.7541
313/313 [=====] - 4s 13ms/step - loss: 0.9102 -
accuracy: 0.7541
test loss: 0.9102030992507935
```

test accuracy: 0.7541000247001648

```
confusion matrix:
[[805  18  20  18  10  17  14  21  43  34]
 [ 15 877   1  17   3   7  13   6  17  44]
 [ 70   4 533  68  97  75  93  45   9   6]
 [ 18   3  30 547  58 210  78  41  11   4]
 [ 20   1  28  67 720  42  55  57   5   5]
 [ 11   1  22 144  40 713  24  42   3   0]
 [  5   2  15  56  26  32 853   6   3   2]
 [  7   0  11  29  37  62  10 831   5   8]
 [ 62  25   6  13   5   6   8   8 841  26]
 [ 23  75   5  14   3  18  12  18  11 821]]
```

[15] :



14

```

model.add(keras.layers.Flatten())
model.add(keras.layers.Dense(64, activation = 'relu'))
model.add(BatchNormalization())
model.add(layers.Dropout(0.2))
model.add(keras.layers.Dense(10, activation = 'softmax'))

model.summary()

```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 32, 32, 16)	448
conv2d_13 (Conv2D)	(None, 32, 32, 16)	2320
batch_normalization_4 (Batch Normalization)	(None, 32, 32, 16)	64
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 16)	0
dropout (Dropout)	(None, 16, 16, 16)	0
conv2d_14 (Conv2D)	(None, 16, 16, 32)	4640
conv2d_15 (Conv2D)	(None, 16, 16, 32)	9248
batch_normalization_5 (Batch Normalization)	(None, 16, 16, 32)	128
max_pooling2d_4 (MaxPooling2D)	(None, 8, 8, 32)	0
dropout_1 (Dropout)	(None, 8, 8, 32)	0
conv2d_16 (Conv2D)	(None, 8, 8, 64)	18496
conv2d_17 (Conv2D)	(None, 8, 8, 64)	36928
batch_normalization_6 (Batch Normalization)	(None, 8, 8, 64)	256
max_pooling2d_5 (MaxPooling2D)	(None, 4, 4, 64)	0

dropout_2 (Dropout)	(None, 4, 4, 64)	0
flatten_2 (Flatten)	(None, 1024)	0
dense_4 (Dense)	(None, 64)	65600
batch_normalization_7 (Batch Normalization)	(None, 64)	256
dropout_3 (Dropout)	(None, 64)	0
dense_5 (Dense)	(None, 10)	650

=====

Total params: 139,034

Trainable params: 138,682

Non-trainable params: 352

-----

```
[17]: model_compile('categorical_crossentropy', 'adam', 32)
      plot(model_history.history)
      test_evaluate(model)
      model_architecture(model)
```

Epoch 1/500

1250/1250 [=====] - 75s 60ms/step - loss: 1.6013 - accuracy: 0.4311 - val\_loss: 1.2272 - val\_accuracy: 0.5618

...

Epoch 17/500

1250/1250 [=====] - 87s 69ms/step - loss: 0.5453 - accuracy: 0.8105 - val\_loss: 0.6458 - val\_accuracy: 0.7771






Figure 7: Loss and validation loss for CNN model with batch normalization, Max pooling and dropout




Figure 8: Accuracy and validation accuracy for CNN model with batch normalization, Max pooling and dropout

```
313/313 [=====] - 4s 14ms/step - loss: 0.6282 -
accuracy: 0.7842
313/313 [=====] - 4s 14ms/step - loss: 0.6282 -
accuracy: 0.7842
test loss: 0.6282299757003784
```

```
test accuracy: 0.7842000126838684
```

```
confusion matrix:
[[891  7  24  9  6  0  2  12  33  16]
 [ 31 854  4  3  1  2  5  1  29  70]
 [ 87  4 696 27 57 48 50 24  6  1]
 [ 42  3  90 544 38 175 60 28 13  7]
 [ 39  2  76  43 717 29 40 48  6  0]
 [ 25  2  76  78  31 733 21 29  5  0]
 [  9  1  47  41  22  18 846 11  5  0]
 [ 15  0  37  18  27  60  2 837  1  3]
 [ 80 10  8  9  2  3  2  3 869 14]
 [ 50 43  5  6  4  1  5  7  24 855]]
```



we already used early stop callback in our models.

## 2 Transfer Learning

Student IDs: 810100462 & 810100410

selected model: 2 -> SqueezeNet

### 2.1 Model Properties

**model architecture:** The SqueezeNet architecture is comprised of “squeeze” and “expand” layers. A squeeze convolutional layer has only  $1 \times 1$  filters. These are fed into an expand layer that has a mix of  $1 \times 1$  and  $3 \times 3$  convolution filters.

The authors of the paper use the term “fire module” to describe a squeeze layer and an expand layer together.

An input image is first sent into a standalone convolutional layer. This layer is followed by 8 “fire modules” which are named “fire2-9”, according to Strategy One above.

Following Strategy Two, the filters per fire module are increased with “simple bypass.” Lastly, SqueezeNet performs max-pooling with a stride of 2 after layers conv1, fire4, fire8, and conv10. According to Strategy Three, pooling is given a relatively late placement, resulting in SqueezeNet with a “complex bypass”.

Below are the details of other parameters used in the network: \* The ReLU activation is applied between all the squeeze and expand layers inside the fire module. \* Dropout layers are added to reduce overfitting, with a probability of 0.5 after the fire9 module. \* There are no fully connected layers used in the network. This design choice was inspired by the Network In Network (NIN) architecture proposed by (Lin et al, 2013). \* SqueezeNet was trained with a learning rate of 0.04, which is linearly decreased throughout the training process. \* The batch size for training is 32, and the network used an Adam Optimizer.

**pros and cons:** SqueezeNet makes the deployment process easier due to its small size.

SqueezeNet has two disadvantages: low classification accuracy and high computational complexity.

**initial preprocess:** All pre-trained models expect input images normalized in the same way, i.e. mini-batches of 3-channel RGB images of shape  $(3 \times H \times W)$ , where H and W are expected to be at least 224. The images have to be loaded in to a range of  $[0, 1]$  and then normalized using mean and std.

### 2.2 Transfer Learning

**Transfer Learning** is a machine learning method where we reuse a pre-trained model as the starting point for a model on a new task.

To put it simply—a model trained on one task is repurposed on a second, related task as an optimization that allows rapid progress when modeling the second task.

By applying transfer learning to a new task, one can achieve significantly higher performance than training with only a small amount of data.

```
[1]: import torch
```

```
[2]: model = torch.hub.load('pytorch/vision:v0.10.0', 'squeezenet1_0',
    ↳ pretrained=True)
    model.eval()
```

Downloading: "https://github.com/pytorch/vision/archive/v0.10.0.zip" to  
/root/.cache/torch/hub/v0.10.0.zip

Downloading: "https://download.pytorch.org/models/squeezenet1\_0-b66bffa10.pth" to  
/root/.cache/torch/hub/checkpoints/squeezenet1\_0-b66bffa10.pth

0%| | 0.00/4.78M [00:00<?, ?B/s]

```
[2]: SqueezeNet(
  (features): Sequential(
    (0): Conv2d(3, 96, kernel_size=(7, 7), stride=(2, 2))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=True)
    (3): Fire(
      (squeeze): Conv2d(96, 16, kernel_size=(1, 1), stride=(1, 1))
      (squeeze_activation): ReLU(inplace=True)
      (expand1x1): Conv2d(16, 64, kernel_size=(1, 1), stride=(1, 1))
      (expand1x1_activation): ReLU(inplace=True)
      (expand3x3): Conv2d(16, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
      (expand3x3_activation): ReLU(inplace=True)
    )
    (4): Fire(
      (squeeze): Conv2d(128, 16, kernel_size=(1, 1), stride=(1, 1))
      (squeeze_activation): ReLU(inplace=True)
      (expand1x1): Conv2d(16, 64, kernel_size=(1, 1), stride=(1, 1))
      (expand1x1_activation): ReLU(inplace=True)
      (expand3x3): Conv2d(16, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
      (expand3x3_activation): ReLU(inplace=True)
    )
    (5): Fire(
      (squeeze): Conv2d(128, 32, kernel_size=(1, 1), stride=(1, 1))
      (squeeze_activation): ReLU(inplace=True)
      (expand1x1): Conv2d(32, 128, kernel_size=(1, 1), stride=(1, 1))
      (expand1x1_activation): ReLU(inplace=True)
      (expand3x3): Conv2d(32, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (expand3x3_activation): ReLU(inplace=True)
    )
    (6): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=True)
    (7): Fire(
      (squeeze): Conv2d(256, 32, kernel_size=(1, 1), stride=(1, 1))
```

```

        (squeeze_activation): ReLU(inplace=True)
        (expand1x1): Conv2d(32, 128, kernel_size=(1, 1), stride=(1, 1))
        (expand1x1_activation): ReLU(inplace=True)
        (expand3x3): Conv2d(32, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        (expand3x3_activation): ReLU(inplace=True)
    )
    (8): Fire(
        (squeeze): Conv2d(256, 48, kernel_size=(1, 1), stride=(1, 1))
        (squeeze_activation): ReLU(inplace=True)
        (expand1x1): Conv2d(48, 192, kernel_size=(1, 1), stride=(1, 1))
        (expand1x1_activation): ReLU(inplace=True)
        (expand3x3): Conv2d(48, 192, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        (expand3x3_activation): ReLU(inplace=True)
    )
    (9): Fire(
        (squeeze): Conv2d(384, 48, kernel_size=(1, 1), stride=(1, 1))
        (squeeze_activation): ReLU(inplace=True)
        (expand1x1): Conv2d(48, 192, kernel_size=(1, 1), stride=(1, 1))
        (expand1x1_activation): ReLU(inplace=True)
        (expand3x3): Conv2d(48, 192, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        (expand3x3_activation): ReLU(inplace=True)
    )
    (10): Fire(
        (squeeze): Conv2d(384, 64, kernel_size=(1, 1), stride=(1, 1))
        (squeeze_activation): ReLU(inplace=True)
        (expand1x1): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1))
        (expand1x1_activation): ReLU(inplace=True)
        (expand3x3): Conv2d(64, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        (expand3x3_activation): ReLU(inplace=True)
    )
    (11): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=True)
    (12): Fire(
        (squeeze): Conv2d(512, 64, kernel_size=(1, 1), stride=(1, 1))
        (squeeze_activation): ReLU(inplace=True)
        (expand1x1): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1))
        (expand1x1_activation): ReLU(inplace=True)
        (expand3x3): Conv2d(64, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        (expand3x3_activation): ReLU(inplace=True)
    )
    )
    (classifier): Sequential(

```

```

(0): Dropout(p=0.5, inplace=False)
(1): Conv2d(512, 1000, kernel_size=(1, 1), stride=(1, 1))
(2): ReLU(inplace=True)
(3): AdaptiveAvgPool2d(output_size=(1, 1))
)
)

```

## 2.3 Implementation

our model can classify animals.

```

[7]: import urllib
url, filename = ("https://github.com/pytorch/hub/raw/master/images/dog.jpg",
    ↪ "dog.jpg")
try: urllib.urlopen().retrieve(url, filename)
except: urllib.request.urlretrieve(url, filename)

```

```

[5]: from PIL import Image
import matplotlib.pyplot as plt
dog = Image.open('dog.jpg')
imgplot = plt.imshow(dog)

```




Figure 10: a custom photo of recognizable categories



```
[8]: from PIL import Image
      from torchvision import transforms

      input_image = Image.open(filename)
      preprocess = transforms.Compose([
          transforms.Resize(299),
          transforms.CenterCrop(299),
          transforms.ToTensor(),
          transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
      ])
      input_tensor = preprocess(input_image)
      input_batch = input_tensor.unsqueeze(0) # create a mini-batch as expected by the
      →model

      # move the input and model to GPU for speed if available
      if torch.cuda.is_available():
          input_batch = input_batch.to('cuda')
          model.to('cuda')

      with torch.no_grad():
          output = model(input_batch)

      probabilities = torch.nn.functional.softmax(output[0], dim=0)
```

```
[9]: !wget https://raw.githubusercontent.com/pytorch/hub/master/imagenet_classes.txt

      # Read the categories
      with open("imagenet_classes.txt", "r") as f:
          categories = [s.strip() for s in f.readlines()]

--2022-04-15 09:09:56--
https://raw.githubusercontent.com/pytorch/hub/master/imagenet_classes.txt
Resolving raw.githubusercontent.com (raw.githubusercontent.com)...
185.199.108.133, 185.199.109.133, 185.199.110.133, ...
Connecting to raw.githubusercontent.com
(raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 10472 (10K) [text/plain]
Saving to: 'imagenet_classes.txt.1'
```

```
imagenet_classes.tx 100%[=====>] 10.23K --.-KB/s in 0s
```

```
2022-04-15 09:09:56 (74.9 MB/s) - 'imagenet_classes.txt.1' saved [10472/10472]
```

```
[10]: _, indices = torch.topk(probabilities, 3)
      for index in indices:
          print('Object {} with probability {}'.format(categories[index],
          ↪probabilities[index]))
```

Object Samoyed with probability 0.8526738882064819

Object Pomeranian with probability 0.03457362577319145

Object West Highland white terrier with probability 0.02477547898888588

### 3 Segmentation

**FCN:** \* remove any dense(fully connected) layers and only use convolution layers. \* downsampling and then do upsampling to get pixel wise prediction to do segmentation. \* also introduce lateral connections, which is combining downsampling feature map and upsampled feature map on same level. \* when combining lateral connections, the paper simply adds the two. \* upsampling is a layer which is initialized as bilinear interpolation but allow the params to be learned.

**UNet:** \* down and upsampling architecture. \* only the final upsampled feature map is utilized. \* designed for segmentation.

import essential libraries:

```
[1]: import tensorflow_datasets as tfds
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras import layers
```

import 'oxford\_iiit\_pet' dataset for training:

```
[2]: dataset, info = tfds.load('oxford_iiit_pet:3.*.*', with_info=True)
```

resize the images and masks to 128x128:

```
[3]: def resize(input_image, input_mask):
    input_image = tf.image.resize(input_image, (128, 128), method="nearest")
    input_mask = tf.image.resize(input_mask, (128, 128), method="nearest")
    return input_image, input_mask
```

create a function to augment the dataset by flipping them horizontally:

```
[4]: def augment(input_image, input_mask):
    if tf.random.uniform(()) > 0.5:
        # Random flipping of the image and mask
        input_image = tf.image.flip_left_right(input_image)
        input_mask = tf.image.flip_left_right(input_mask)
    return input_image, input_mask
```

normalize the dataset by scaling the images to the range of [-1, 1] and decreasing the image mask by 1:

```
[5]: def normalize(input_image, input_mask):
    input_image = tf.cast(input_image, tf.float32) / 255.0
    input_mask -= 1
    return input_image, input_mask
```

functions to preprocess the training and test datasets:

```
[6]: def load_image_train(datapoint):
    input_image = datapoint["image"]
    input_mask = datapoint["segmentation_mask"]
    input_image, input_mask = resize(input_image, input_mask)
    input_image, input_mask = augment(input_image, input_mask)
    input_image, input_mask = normalize(input_image, input_mask)
    return input_image, input_mask
def load_image_test(datapoint):
    input_image = datapoint["image"]
    input_mask = datapoint["segmentation_mask"]
    input_image, input_mask = resize(input_image, input_mask)
    input_image, input_mask = normalize(input_image, input_mask)
    return input_image, input_mask
```

build an input pipeline:

```
[7]: train_dataset = dataset["train"].map(load_image_train, num_parallel_calls=tf.
    ↳data.AUTOTUNE)
test_dataset = dataset["test"].map(load_image_test, num_parallel_calls=tf.data.
    ↳AUTOTUNE)
```

making dataset ready for training:

```
[8]: BATCH_SIZE = 64
BUFFER_SIZE = 1000
train_batches = train_dataset.cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE).
    ↳repeat()
train_batches = train_batches.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
validation_batches = test_dataset.take(3000).batch(BATCH_SIZE)
test_batches = test_dataset.skip(3000).take(669).batch(BATCH_SIZE)
```

visualize a random sample image and its mask from the training dataset:

```
[14]: def display(display_list):
    plt.figure(figsize=(15, 15))
    title = ["Input Image", "True Mask", "Predicted Mask"]
    for i in range(len(display_list)):
        plt.subplot(1, len(display_list), i+1)
        plt.title(title[i])
        plt.imshow(tf.keras.utils.array_to_img(display_list[i]))
        plt.axis("off")
    plt.show()
sample_batch = next(iter(train_batches))
random_index = np.random.choice(sample_batch[0].shape[0])
sample_image, sample_mask = sample_batch[0][random_index],
    ↳sample_batch[1][random_index]
display([sample_image, sample_mask])
```




Figure 11: Sample input image and mask

building blocks:

```
[15]: def double_conv_block(x, n_filters):
    # Conv2D then ReLU activation
    x = layers.Conv2D(n_filters, 3, padding = "same", activation = "relu",
    ↪kernel_initializer = "he_normal")(x)
    # Conv2D then ReLU activation
    x = layers.Conv2D(n_filters, 3, padding = "same", activation = "relu",
    ↪kernel_initializer = "he_normal")(x)
    return x
```

```
[16]: def downsample_block(x, n_filters):
    f = double_conv_block(x, n_filters)
    p = layers.MaxPool2D(2)(f)
    p = layers.Dropout(0.3)(p)
    return f, p
```

```
[17]: def upsample_block(x, conv_features, n_filters):
    # upsample
    x = layers.Conv2DTranspose(n_filters, 3, 2, padding="same")(x)
    # concatenate
    x = layers.concatenate([x, conv_features])
    # dropout
    x = layers.Dropout(0.3)(x)
    # Conv2D twice with ReLU activation
```

```
x = double_conv_block(x, n_filters)
return x
```

Unet model:

```
[13]: def build_unet_model():
    # inputs
    inputs = layers.Input(shape=(128,128,3))

    # encoder: contracting path - downsample
    # 1 - downsample
    f1, p1 = downsample_block(inputs, 64)
    # 2 - downsample
    f2, p2 = downsample_block(p1, 128)
    # 3 - downsample
    f3, p3 = downsample_block(p2, 256)
    # 4 - downsample
    f4, p4 = downsample_block(p3, 512)

    # 5 - bottleneck
    bottleneck = double_conv_block(p4, 1024)

    # decoder: expanding path - upsample
    # 6 - upsample
    u6 = upsample_block(bottleneck, f4, 512)
    # 7 - upsample
    u7 = upsample_block(u6, f3, 256)
    # 8 - upsample
    u8 = upsample_block(u7, f2, 128)
    # 9 - upsample
    u9 = upsample_block(u8, f1, 64)

    # outputs
    outputs = layers.Conv2D(3, 1, padding="same", activation = "softmax")(u9)

    # unet model with Keras Functional API
    unet_model = tf.keras.Model(inputs, outputs, name="U-Net")

    return unet_model
```

```
[14]: unet_model = build_unet_model()
```

compile and train Unet:

```
[15]: unet_model.compile(optimizer=tf.keras.optimizers.Adam(),
                        loss="sparse_categorical_crossentropy",
                        metrics="accuracy")
```

```
[16]: NUM_EPOCHS = 10

TRAIN_LENGTH = info.splits["train"].num_examples
STEPS_PER_EPOCH = TRAIN_LENGTH // BATCH_SIZE

VAL_SUBSPLITS = 5
TEST_LENGTH = info.splits["test"].num_examples
VALIDATION_STEPS = TEST_LENGTH // BATCH_SIZE // VAL_SUBSPLITS

model_history = unet_model.fit(train_batches,
                               epochs=NUM_EPOCHS,
                               steps_per_epoch=STEPS_PER_EPOCH,
                               validation_steps=VALIDATION_STEPS,
                               validation_data=test_batches)

Epoch 1/10
57/57 [=====] - 162s 2s/step - loss: 2.1894 - accuracy:
0.5593 - val_loss: 0.9181 - val_accuracy: 0.5864
Epoch 2/10
57/57 [=====] - 130s 2s/step - loss: 0.8967 - accuracy:
0.6057 - val_loss: 0.8755 - val_accuracy: 0.5864
Epoch 3/10
57/57 [=====] - 120s 2s/step - loss: 0.7823 - accuracy:
0.6693 - val_loss: 0.7104 - val_accuracy: 0.7190
Epoch 4/10
57/57 [=====] - 121s 2s/step - loss: 0.6857 - accuracy:
0.7217 - val_loss: 0.6645 - val_accuracy: 0.7264
Epoch 5/10
57/57 [=====] - 121s 2s/step - loss: 0.6481 - accuracy:
0.7387 - val_loss: 0.6407 - val_accuracy: 0.7337
Epoch 6/10
57/57 [=====] - 120s 2s/step - loss: 0.6147 - accuracy:
0.7520 - val_loss: 0.6235 - val_accuracy: 0.7424
Epoch 7/10
57/57 [=====] - 121s 2s/step - loss: 0.5955 - accuracy:
0.7619 - val_loss: 0.6218 - val_accuracy: 0.7432
Epoch 8/10
57/57 [=====] - 121s 2s/step - loss: 0.5597 - accuracy:
0.7775 - val_loss: 0.5804 - val_accuracy: 0.7636
Epoch 9/10
57/57 [=====] - 121s 2s/step - loss: 0.5139 - accuracy:
0.7976 - val_loss: 0.4846 - val_accuracy: 0.8100
Epoch 10/10
57/57 [=====] - 121s 2s/step - loss: 0.4784 - accuracy:
0.8130 - val_loss: 0.4545 - val_accuracy: 0.8228
```

Unet prediction:

```
[54]: def create_mask(pred_mask):
    pred_mask = tf.argmax(pred_mask, axis=-1)
    pred_mask = pred_mask[..., tf.newaxis]
    return pred_mask[0]

def show_predictions_unet(dataset=None, num=1):
    if dataset:
        for image, mask in dataset.take(num):
            pred_mask = unet_model.predict(image)
            display([image[0], mask[0], create_mask(pred_mask)])
    else:
        display([sample_image, sample_mask,
                  create_mask(model.predict(sample_image[tf.newaxis, ...]))])

count = 0
for i in test_batches:
    count +=1
print("number of batches:", count)
```

number of batches: 11

```
[18]: show_predictions_unet(test_batches)
```



Figure 12: U-Net model prediction on test sample

fcn model:

```
[48]: def build_fcn_model():
    # inputs
    inputs = layers.Input(shape=(128,128,3))

    # encoder: contracting path - downsample
```



```

# 1 - downsample
f1, p1 = downsample_block(inputs, 128)
# 2 - downsample
f2, p2 = downsample_block(p1, 128)
# 3 - downsample
f3, p3 = downsample_block(p2, 128)
# 4 - downsample
f4, p4 = downsample_block(p3, 128)

# 5 - bottleneck
bottleneck = double_conv_block(p4, 512)

# decoder: expanding path - upsample
# 6 - upsample
u6 = upsample_block(bottleneck, f4, 128)
# 7 - upsample
u7 = upsample_block(u6, f3, 128)
# 8 - upsample
u8 = upsample_block(u7, f2, 128)
# 9 - upsample
u9 = upsample_block(u8, f1, 128)

# outputs
outputs = layers.Conv2D(3, 1, padding="same", activation = "softmax")(u9)

# fcn model with Keras Functional API
fcn_model = tf.keras.Model(inputs, outputs, name="FCN")

return fcn_model

```

```
[49]: fcn_model = build_fcn_model()
```

compile and train fcn model:

```
[50]: fcn_model.compile(optimizer=tf.keras.optimizers.Adam(),
                        loss="sparse_categorical_crossentropy",
                        metrics="accuracy")
```

```
[51]: NUM_EPOCHS = 10

TRAIN_LENGTH = info.splits["train"].num_examples
STEPS_PER_EPOCH = TRAIN_LENGTH // BATCH_SIZE

VAL_SUBSPLITS = 5
TEST_LENGTH = info.splits["test"].num_examples
VALIDATION_STEPS = TEST_LENGTH // BATCH_SIZE // VAL_SUBSPLITS

```

```

model_history = fcn_model.fit(train_batches,
                              epochs=NUM_EPOCHS,
                              steps_per_epoch=STEPS_PER_EPOCH,
                              validation_steps=VALIDATION_STEPS,
                              validation_data=test_batches)

```

```

Epoch 1/10
57/57 [=====] - 278s 5s/step - loss: 0.9992 - accuracy:
0.5656 - val_loss: 0.8899 - val_accuracy: 0.5864
Epoch 2/10
57/57 [=====] - 180s 3s/step - loss: 0.8151 - accuracy:
0.6151 - val_loss: 0.7467 - val_accuracy: 0.6691
Epoch 3/10
57/57 [=====] - 173s 3s/step - loss: 0.6820 - accuracy:
0.7131 - val_loss: 0.6134 - val_accuracy: 0.7567
Epoch 4/10
57/57 [=====] - 173s 3s/step - loss: 0.6216 - accuracy:
0.7513 - val_loss: 0.5687 - val_accuracy: 0.7751
Epoch 5/10
57/57 [=====] - 173s 3s/step - loss: 0.5564 - accuracy:
0.7818 - val_loss: 0.5354 - val_accuracy: 0.7920
Epoch 6/10
57/57 [=====] - 173s 3s/step - loss: 0.5262 - accuracy:
0.7964 - val_loss: 0.5607 - val_accuracy: 0.7842
Epoch 7/10
57/57 [=====] - 173s 3s/step - loss: 0.4887 - accuracy:
0.8114 - val_loss: 0.4386 - val_accuracy: 0.8287
Epoch 8/10
57/57 [=====] - 173s 3s/step - loss: 0.4328 - accuracy:
0.8335 - val_loss: 0.4242 - val_accuracy: 0.8349
Epoch 9/10
57/57 [=====] - 173s 3s/step - loss: 0.4170 - accuracy:
0.8399 - val_loss: 0.3844 - val_accuracy: 0.8511
Epoch 10/10
57/57 [=====] - 172s 3s/step - loss: 0.4100 - accuracy:
0.8420 - val_loss: 0.3853 - val_accuracy: 0.8527

```

prediction by fcn:

```

[52]: def show_predictions_fcn(dataset=None, num=1):
      if dataset:
          for image, mask in dataset.take(num):
              pred_mask = fcn_model.predict(image)
              display([image[0], mask[0], create_mask(pred_mask)])
      else:
          display([sample_image, sample_mask,
                   create_mask(model.predict(sample_image[tf.newaxis, ...]))])

```

```
[55]: show_predictions_fcn(test_batches)
```



Figure 13: FCN model prediction on test sample

the accuracy and validation accuracy for fcn model is greater than model for nearly 3 percent.

## 4 Object Detection

### 4.1 YOLO v1 and v2

The changes from YOLO to YOLO v2:

- **Batch Normalization:** YOLO V2 normalise the input layer by altering slightly and scaling the activations. Batch normalization decreases the shift in unit value in the hidden layer and by doing so it improves the stability of the neural network. By adding batch normalization to convolutional layers in the architecture MAP (mean average precision) has been improved by 2%. It also helped the model regularise and overfitting has been reduced overall.
- **Higher Resolution Classifier:** the input size in YOLO v2 has been increased from 224x224 to 448x448. The increase in the input size of the image has improved the MAP (mean average precision) upto 4%. This increase in input size is been applied while training the YOLO v2 architecture DarkNet 19 on ImageNet dataset.
- **Anchor Boxes:** one of the most notable changes which can visible in YOLO v2 is the introduction the anchor boxes. YOLO v2 does classification and prediction in a single framework. These anchor boxes are responsible for predicting bounding box and this anchor boxes are designed for a given dataset by using clustering(k-means clustering).
- **Fine-Grained Features:** one of the main issued that has to be addressed in the YOLO v1 is that detection of smaller objects on the image. This has been resolved in the YOLO v2 divides the image into 13x13 grid cells which is smaller when compared to its previous version. This enables the yolo v2 to identify or localize the smaller objects in the image and also effective with the larger objects.
- **Multi-Scale Training:** on YOLO v1 has a weakness detecting objects with different input sizes which says that if YOLO is trained with small images of a particular object it has issues detecting the same object on image of bigger size. This has been resolved to a great extent in YOLO v2 where it is trained with random images with different dimensions range between 320x320 to 608x608. This allows the network to learn and predict the objects from various input dimensions with accuracy.
- **Darknet 19:** YOLO v2 uses Darknet 19 architecture with 19 convolutional layers and 5 max pooling layers and a softmax layer for classification objects. Darknet is a neural network framework written in C language and CUDA. It's really fast in object detection which is very important for predicting in real-time.

With the advancements in several categories in YOLO v2 is better, faster, and stronger. With Multi-Scale Training now the network is able to detect and classify objects with different configurations and dimensions. YOLO v2 has seen a great improvement in detecting smaller objects with much more accuracy which it lacked in its predecessor version.

### 4.2 YOLO 4 & YOLO 5:

As a modified version of YOLOv3, YOLO4. uses Cross Stage Partial Network (CSPNet) in Darknet, creating a new feature extractor backbone called CSPDarknet53. The convolution architecture is based on modified DenseNet. It transfers a copy of feature map from the base layer to the next layer through dense block. The advantages of using DenseNet include the diminishing gradient vanishing problems, boosting backpropagation, removal of the computational bottleneck,

and improved learning. Neck is composed of spatial pyramid pooling (SPP) layer and PANet path aggregation. SPP layer and PANet path aggregation are used for feature aggregation to improve the receptive field and short out important features from the backbone. In addition, the head is composed of YOLO layer. First, the image is fed to CSPDarknet53 for feature extraction and then fed to path aggregation network PANet for fusion. Finally, YOLO layer generates the results, similar to YOLOv3 YOLOv4 uses bag of freebies and bag of specials to improve the algorithm performance. Bag of freebies includes Complete IOU loss (CIOU), drop block regularization and different augmentation techniques. Bags of specials includes mish activation, Diou-NMS and modified the path aggregation networks.

However, YOLOv5 is different from the previous releases. It utilizes PyTorch instead of Darknet. It utilizes CSPDarknet53 as backbone. This backbone solves the repetitive gradient information in large backbones and integrates gradient change into feature map that reduces the inference speed, increases accuracy, and reduces the model size by decreasing the parameters. It uses path aggregation network (PANet) as neck to boost the information flow. PANet adopts a new feature pyramid network (FPN) that includes several bottom ups and top down layers. This improves the propagation of low level features in the model. PANet improves the localization in lower layers, which enhances the localization accuracy of the object. In addition, the head in YOLOv5 is the same as YOLOv4 and YOLOv3 which generates three different output of feature maps to achieve multi scale prediction. It also helps to enhance the prediction of small to large objects efficiently in the model. The image is fed to CSPDarknet53 for feature extraction and again fed to PANet for feature fusion. Finally, the YOLO layer generates the results.

### 4.3 Two-stage vs One-stage Detectors:

In two-stage Detectors, first, the model proposes a set of regions of interests by select search or regional proposal network. The proposed regions are sparse as the potential bounding box candidates can be infinite. Then a classifier only processes the region candidates. But in one-stage Detectors Single convolutional network predicts the bounding boxes and the class probabilities for these boxes.

**Example of two-stage Detectors:** R-CNN family

**Examples of one-stage Detectors:** YOLO & SSD

### 4.4 YOLO v5 implementation:

Install Dependencies:

```
[ ]: # clone YOLOv5 repository
!git clone https://github.com/ultralytics/yolov5 # clone repo
%cd yolov5
!git reset --hard 886f1c03d839575afecb059acbf74296fad395b6
```

```
fatal: destination path 'yolov5' already exists and is not an empty directory.
/content/yolov5
HEAD is now at 886f1c0 DDP after autoanchor reorder (#2421)
```

```
[ ]: # install dependencies as necessary
!pip install -qr requirements.txt # install dependencies (ignore errors)
import torch

from IPython.display import Image, clear_output # to display images
from utils.google_utils import gdrive_download # to download models/datasets

# clear_output()
print('Setup complete. Using torch %s %s' % (torch.__version__, torch.cuda.
→get_device_properties(0) if torch.cuda.is_available() else 'CPU'))
```

Setup complete. Using torch 1.10.0+cu111 \_CudaDeviceProperties(name='Tesla K80', major=3, minor=7, total\_memory=11441MB, multi\_processor\_count=13)

```
[ ]: %cd /content
!curl -L "https://app.roboflow.com/ds/F2Eo4hPmah?key=8ZdNTyoAP8" > roboflow.zip;
→unzip roboflow.zip; rm roboflow.zip
```

```
/content
% Total      % Received % Xferd  Average Speed   Time    Time       Time  Current
           Dload  Upload   Total     Spent    Left     Speed
100  896  100  896    0     0    746      0  0:00:01  0:00:01  --:--:--   746
100 5698k  100 5698k    0     0 2279k      0  0:00:02  0:00:02  --:--:--  245M
Archive:  roboflow.zip
  extracting: README.roboflow.txt
  extracting: data.yaml
    creating: test/
    creating: test/images/
  extracting: test/images/cap_100_jpg.rf.04d2e5f139a204fe67ba4a6e318b9352.jpg
...
  extracting: valid/labels/cap_91_jpg.rf.f4c440ea8c8a9f0304073f9b6336bba2.txt
```

```
[ ]: %cat data.yaml
```

```
train: ../train/images
val: ../valid/images

nc: 1
names: ['Caps']
```

```
[ ]: import yaml
with open("data.yaml", 'r') as stream:
    num_classes = str(yaml.safe_load(stream)['nc'])
```

```
[ ]: %cat /content/yolov5/models/yolov5s.yaml
```

```
# parameters
nc: 80 # number of classes
```

```

depth_multiple: 0.33 # model depth multiple
width_multiple: 0.50 # layer channel multiple

# anchors
anchors:
  - [10,13, 16,30, 33,23] # P3/8
  - [30,61, 62,45, 59,119] # P4/16
  - [116,90, 156,198, 373,326] # P5/32

# YOLOv5 backbone
backbone:
  # [from, number, module, args]
  [[-1, 1, Focus, [64, 3]], # 0-P1/2
   [-1, 1, Conv, [128, 3, 2]], # 1-P2/4
   [-1, 3, C3, [128]],
   [-1, 1, Conv, [256, 3, 2]], # 3-P3/8
   [-1, 9, C3, [256]],
   [-1, 1, Conv, [512, 3, 2]], # 5-P4/16
   [-1, 9, C3, [512]],
   [-1, 1, Conv, [1024, 3, 2]], # 7-P5/32
   [-1, 1, SPP, [1024, [5, 9, 13]]],
   [-1, 3, C3, [1024, False]], # 9
  ]

# YOLOv5 head
head:
  [[-1, 1, Conv, [512, 1, 1]],
   [-1, 1, nn.Upsample, [None, 2, 'nearest']],
   [[-1, 6], 1, Concat, [1]], # cat backbone P4
   [-1, 3, C3, [512, False]], # 13

   [-1, 1, Conv, [256, 1, 1]],
   [-1, 1, nn.Upsample, [None, 2, 'nearest']],
   [[-1, 4], 1, Concat, [1]], # cat backbone P3
   [-1, 3, C3, [256, False]], # 17 (P3/8-small)

   [-1, 1, Conv, [256, 3, 2]],
   [[-1, 14], 1, Concat, [1]], # cat head P4
   [-1, 3, C3, [512, False]], # 20 (P4/16-medium)

   [-1, 1, Conv, [512, 3, 2]],
   [[-1, 10], 1, Concat, [1]], # cat head P5
   [-1, 3, C3, [1024, False]], # 23 (P5/32-large)

  [[17, 20, 23], 1, Detect, [nc, anchors]], # Detect(P3, P4, P5)
  ]

```

```
[ ]: from IPython.core.magic import register_line_cell_magic
```

```
@register_line_cell_magic
def writetemplate(line, cell):
    with open(line, 'w') as f:
        f.write(cell.format(**globals()))
```

```
[ ]: %%writetemplate /content/yolov5/models/custom_yolov5s.yaml
```

```
# parameters
nc: {num_classes} # number of classes
depth_multiple: 0.33 # model depth multiple
width_multiple: 0.50 # layer channel multiple

# anchors
anchors:
  - [10,13, 16,30, 33,23] # P3/8
  - [30,61, 62,45, 59,119] # P4/16
  - [116,90, 156,198, 373,326] # P5/32

# YOLOv5 backbone
backbone:
  # [from, number, module, args]
  [[-1, 1, Focus, [64, 3]], # 0-P1/2
  [-1, 1, Conv, [128, 3, 2]], # 1-P2/4
  [-1, 3, BottleneckCSP, [128]],
  [-1, 1, Conv, [256, 3, 2]], # 3-P3/8
  [-1, 9, BottleneckCSP, [256]],
  [-1, 1, Conv, [512, 3, 2]], # 5-P4/16
  [-1, 9, BottleneckCSP, [512]],
  [-1, 1, Conv, [1024, 3, 2]], # 7-P5/32
  [-1, 1, SPP, [1024, [5, 9, 13]]],
  [-1, 3, BottleneckCSP, [1024, False]], # 9
  ]

# YOLOv5 head
head:
  [[-1, 1, Conv, [512, 1, 1]],
  [-1, 1, nn.Upsample, [None, 2, 'nearest']],
  [[-1, 6], 1, Concat, [1]], # cat backbone P4
  [-1, 3, BottleneckCSP, [512, False]], # 13

  [-1, 1, Conv, [256, 1, 1]],
  [-1, 1, nn.Upsample, [None, 2, 'nearest']],
  [[-1, 4], 1, Concat, [1]], # cat backbone P3
  [-1, 3, BottleneckCSP, [256, False]], # 17 (P3/8-small)
```



```

[-1, 1, Conv, [256, 3, 2]],
[[-1, 14], 1, Concat, [1]], # cat head P4
[-1, 3, BottleneckCSP, [512, False]], # 20 (P4/16-medium)

[-1, 1, Conv, [512, 3, 2]],
[[-1, 10], 1, Concat, [1]], # cat head P5
[-1, 3, BottleneckCSP, [1024, False]], # 23 (P5/32-large)

[[17, 20, 23], 1, Detect, [nc, anchors]], # Detect(P3, P4, P5)
]

```

### Train Custom YOLOv5 Detector:

```

[ ]: %%time
%cd /content/yolov5/
!python train.py --img 416 --batch 16 --epochs 100 --data '../data.yaml' --cfg ./
    models/custom_yolov5s.yaml --weights '' --name yolov5s_results --cache

```

/content/yolov5

github: WARNING: code is out of date by 1006 commits. Use 'git pull' to update or 'git clone https://github.com/ultralytics/yolov5' to download latest.

YOLOv5 v4.0-126-g886f1c0 torch 1.10.0+cu111 CUDA:0 (Tesla K80, 11441.1875MB)

Namespace(adam=False, batch\_size=16, bucket='', cache\_images=True, cfg='./models/custom\_yolov5s.yaml', data='../data.yaml', device='', entity=None, epochs=100, evolve=False, exist\_ok=False, global\_rank=-1, hyp='data/hyp.scratch.yaml', image\_weights=False, img\_size=[416, 416], linear\_lr=False, local\_rank=-1, log\_artifacts=False, log\_imgs=16, multi\_scale=False, name='yolov5s\_results', noautoanchor=False, nosave=False, notest=False, project='runs/train', quad=False, rect=False, resume=False, save\_dir='runs/train/yolov5s\_results', single\_cls=False, sync\_bn=False, total\_batch\_size=16, weights='', workers=8, world\_size=1)

wandb: Install Weights & Biases for YOLOv5 logging with 'pip install wandb' (recommended)

Start Tensorboard with "tensorboard --logdir runs/train", view at http://localhost:6006/

hyperparameters: lr0=0.01, lrf=0.2, momentum=0.937, weight\_decay=0.0005, warmup\_epochs=3.0, warmup\_momentum=0.8, warmup\_bias\_lr=0.1, box=0.05, cls=0.5, cls\_pw=1.0, obj=1.0, obj\_pw=1.0, iou\_t=0.2, anchor\_t=4.0, fl\_gamma=0.0, hsv\_h=0.015, hsv\_s=0.7, hsv\_v=0.4, degrees=0.0, translate=0.1, scale=0.5, shear=0.0, perspective=0.0, flipud=0.0, fliplr=0.5, mosaic=1.0, mixup=0.0

	from	n	params	module	
arguments					
0	-1	1	3520	models.common.Focus	[3,
32, 3]					

1	-1 1	18560	models.common.Conv	[32,
64, 3, 2]				
2	-1 1	19904	models.common.BottleneckCSP	[64,
64, 1]				
3	-1 1	73984	models.common.Conv	[64,
128, 3, 2]				
4	-1 1	161152	models.common.BottleneckCSP	
[128, 128, 3]				
5	-1 1	295424	models.common.Conv	
[128, 256, 3, 2]				
6	-1 1	641792	models.common.BottleneckCSP	
[256, 256, 3]				
7	-1 1	1180672	models.common.Conv	
[256, 512, 3, 2]				
8	-1 1	656896	models.common.SPP	
[512, 512, [5, 9, 13]]				
9	-1 1	1248768	models.common.BottleneckCSP	
[512, 512, 1, False]				
10	-1 1	131584	models.common.Conv	
[512, 256, 1, 1]				
11	-1 1	0	torch.nn.modules.upsampling.Upsample	
[None, 2, 'nearest']				
12	[-1, 6] 1	0	models.common.Concat	[1]
13	-1 1	378624	models.common.BottleneckCSP	
[512, 256, 1, False]				
14	-1 1	33024	models.common.Conv	
[256, 128, 1, 1]				
15	-1 1	0	torch.nn.modules.upsampling.Upsample	
[None, 2, 'nearest']				
16	[-1, 4] 1	0	models.common.Concat	[1]
17	-1 1	95104	models.common.BottleneckCSP	
[256, 128, 1, False]				
18	-1 1	147712	models.common.Conv	
[128, 128, 3, 2]				
19	[-1, 14] 1	0	models.common.Concat	[1]
20	-1 1	313088	models.common.BottleneckCSP	
[256, 256, 1, False]				
21	-1 1	590336	models.common.Conv	
[256, 256, 3, 2]				
22	[-1, 10] 1	0	models.common.Concat	[1]
23	-1 1	1248768	models.common.BottleneckCSP	
[512, 512, 1, False]				
24	[17, 20, 23] 1	16182	models.yolo.Detect	[1,
[[10, 13, 16, 30, 33, 23], [30, 61, 62, 45, 59, 119], [116, 90, 156, 198, 373, 326]], [128, 256, 512]]				

/usr/local/lib/python3.7/dist-packages/torch/functional.py:445: UserWarning:  
torch.meshgrid: in an upcoming release, it will be required to pass the indexing  
argument. (Triggered internally at

```

../aten/src/ATen/native/TensorShape.cpp:2157.)
  return _VF.meshgrid(tensors, **kwargs) # type: ignore[attr-defined]
Model Summary: 283 layers, 7255094 parameters, 7255094 gradients, 16.8 GFLOPS

Scaled weight_decay = 0.0005
Optimizer groups: 62 .bias, 70 conv.weight, 59 other
train: Scanning '../train/labels' for images and labels... 140
found, 0 missing, 0 empty, 0 corrupted: 100% 140/140 [00:00<00:00, 2087.96it/s]
train: New cache created: ../train/labels.cache
train: Caching images (0.1GB): 100% 140/140 [00:00<00:00,
429.69it/s]
val: Scanning '../valid/labels' for images and labels... 40 found,
0 missing, 0 empty, 0 corrupted: 100% 40/40 [00:00<00:00, 1733.99it/s]
val: New cache created: ../valid/labels.cache
val: Caching images (0.0GB): 100% 40/40 [00:00<00:00, 283.62it/s]
Plotting labels...

autoanchor: Analyzing anchors... anchors/target = 5.36, Best
Possible Recall (BPR) = 0.9931
Image sizes 416 train, 416 test
Using 2 dataloader workers
Logging results to runs/train/yolov5s_results
Starting training for 100 epochs...

   Epoch  gpu_mem    box    obj    cls  total  targets  img_size
   0/99    1.79G    0.1086  0.0208     0   0.1294      26
416: 100% 9/9 [00:08<00:00, 1.06it/s]
          Class    Images  Targets          P          R      mAP@.5
mAP@.5:.95: 100% 2/2 [00:02<00:00, 1.40s/it]
          all        40        41      0.0051      0.976      0.00565
0.000752

...

   Epoch  gpu_mem    box    obj    cls  total  targets  img_size
   99/99   1.82G    0.03599  0.01685     0   0.05284      15
416: 100% 9/9 [00:03<00:00, 2.36it/s]
          Class    Images  Targets          P          R      mAP@.5
mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.63it/s]
          all        40        41      0.769      0.975      0.818
0.448
Optimizer stripped from runs/train/yolov5s_results/weights/last.pt, 14.7MB
Optimizer stripped from runs/train/yolov5s_results/weights/best.pt, 14.7MB
100 epochs completed in 0.143 hours.

CPU times: user 5.96 s, sys: 882 ms, total: 6.84 s
Wall time: 8min 56s

```

Evaluate Custom YOLOv5 Detector Performance:

```
[ ]: %load_ext tensorboard
      %tensorboard --logdir runs
```

<IPython.core.display.Javascript object>

```
[ ]: from utils.plots import plot_results # plot results.txt as results.png
      Image(filename='/content/yolov5/runs/train/yolov5s_results/results.png',
            ↪width=1000) # view results.png
```

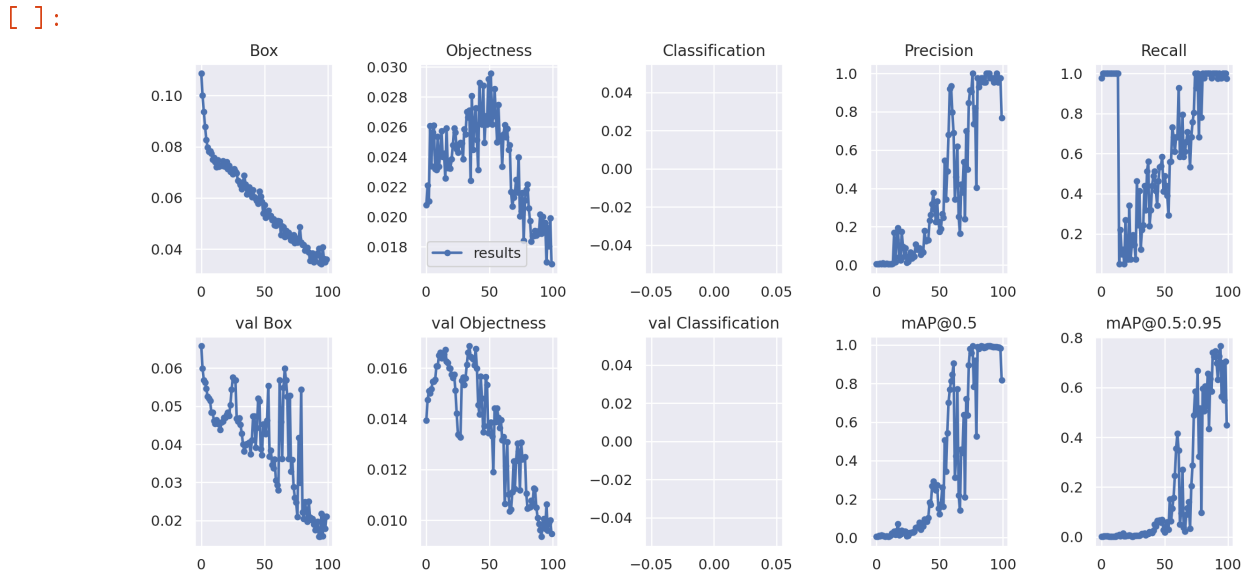


Figure 14: YOLO v5 detector evaluation

Visualize Our Training Data with Labels:

```
[ ]: # first, display our ground truth data
      print("GROUND TRUTH TRAINING DATA:")
      Image(filename='/content/yolov5/runs/train/yolov5s_results/test_batch0_labels.
            ↪jpg', width=900)
```

GROUND TRUTH TRAINING DATA:

```
[ ]:
```



Figure 15: Ground truth training data samples with labels

```
[ ]: # print out an augmented training example
print("GROUND TRUTH AUGMENTED TRAINING DATA:")
Image(filename='/content/yolov5/runs/train/yolov5s_results/train_batch0.jpg',
      width=900)
```

GROUND TRUTH AUGMENTED TRAINING DATA:

```
[ ]:
```






Figure 16: Augmented ground truth training data samples with labels

```
[ ]: %ls runs/
```

```
train/
```

```
[ ]: %ls runs/train/yolov5s_results/weights
```

```
best.pt last.pt
```

```
[ ]: %cd /content/yolov5/  
!python detect.py --weights runs/train/yolov5s_results/weights/best.pt --img 416  
→--conf 0.4 --source ../test/images
```

/content/yolov5

```
Namespace(agnostic_nms=False, augment=False, classes=None, conf_thres=0.4,  
device='', exist_ok=False, img_size=416, iou_thres=0.45, name='exp',  
project='runs/detect', save_conf=False, save_txt=False, source='../test/images',  
update=False, view_img=False,  
weights=['runs/train/yolov5s_results/weights/best.pt'])  
YOL0v5 v4.0-126-g886f1c0 torch 1.10.0+cu111 CUDA:0 (Tesla K80, 11441.1875MB)
```

Fusing layers...

```
/usr/local/lib/python3.7/dist-packages/torch/functional.py:445: UserWarning:  
torch.meshgrid: in an upcoming release, it will be required to pass the indexing  
argument. (Triggered internally at  
../aten/src/ATen/native/TensorShape.cpp:2157.)
```

```
    return _VF.meshgrid(tensors, **kwargs) # type: ignore[attr-defined]
```

Model Summary: 232 layers, 7246518 parameters, 0 gradients, 16.8 GFLOPS

image 1/20 /content/yolov5/./test/images/cap\_100\_jpg.rf.04d2e5f139a204fe67ba4a6e318b9352.jpg: 416x416 1 Caps, Done. (0.033s)

image 2/20 /content/yolov5/./test/images/cap\_101\_jpg.rf.bcafe407a27e06e31019d2802a5f267a.jpg: 416x416 1 Caps, Done. (0.033s)

image 3/20 /content/yolov5/./test/images/cap\_102\_jpg.rf.67265b95ec5339ff91158ab2ea35e9ed.jpg: 416x416 1 Caps, Done. (0.033s)

image 4/20 /content/yolov5/./test/images/cap\_113\_jpg.rf.85dd06988ea675614c4bb3698937e9e6.jpg: 416x416 1 Caps, Done. (0.032s)

image 5/20 /content/yolov5/./test/images/cap\_121\_jpg.rf.99cc4b1767233b440b7397bfe1f77bdf.jpg: 416x416 1 Caps, Done. (0.033s)

image 6/20 /content/yolov5/./test/images/cap\_125\_jpg.rf.1ba0709ce7f0a493a741cfd b1c2f996d.jpg: 416x416 1 Caps, Done. (0.032s)

image 7/20 /content/yolov5/./test/images/cap\_12\_jpg.rf.3188210d6200a6602a3a0cef13a85418.jpg: 416x416 1 Caps, Done. (0.033s)

image 8/20 /content/yolov5/./test/images/cap\_14\_jpg.rf.6b8857c5e479723cd280dedf8a98ad0b.jpg: 416x416 1 Caps, Done. (0.033s)

image 9/20 /content/yolov5/./test/images/cap\_151\_jpg.rf.43f8a7c779e09a08879da99fd13ac9cc.jpg: 416x416 1 Caps, Done. (0.032s)

image 10/20 /content/yolov5/./test/images/cap\_154\_jpg.rf.604207823d49f08e5b013fffa6ff2110.jpg: 416x416 1 Caps, Done. (0.032s)

image 11/20 /content/yolov5/./test/images/cap\_165\_jpg.rf.dd42c9c02ac812af1e415f0f31119a2a.jpg: 416x416 1 Caps, Done. (0.033s)

image 12/20 /content/yolov5/./test/images/cap\_189\_jpg.rf.000b2581f1304c8f69805fe09986ebcb.jpg: 416x416 1 Caps, Done. (0.033s)

image 13/20 /content/yolov5/./test/images/cap\_25\_jpg.rf.5055e576f507088c0986dc5765961294.jpg: 416x416 1 Caps, Done. (0.032s)


image 14/20 /content/yolov5/./test/images/cap\_53\_jpg.rf.03389dd7ec7c331a6fe6eaf2e0adbb4a.jpg: 416x416 1 Caps, Done. (0.032s)

```
image 15/20 /content/yolov5/./test/images/cap_68_jpg.rf.781215a4e29d4f6d03858e3f983e7fb8.jpg: 416x416 Done. (0.032s)
image 16/20 /content/yolov5/./test/images/cap_72_jpg.rf.e2c94bdfc62bd0895a30066fb847ede4.jpg: 416x416 1 Caps, Done. (0.032s)
image 17/20 /content/yolov5/./test/images/cap_79_jpg.rf.e6b2975dc8af5a4724007d5161342f6a.jpg: 416x416 1 Caps, Done. (0.032s)
image 18/20 /content/yolov5/./test/images/cap_83_jpg.rf.d337aa4233ab8ffcce1668a55bb761ab.jpg: 416x416 1 Caps, Done. (0.031s)
image 19/20 /content/yolov5/./test/images/cap_93_jpg.rf.3b41aa7c94cab22004cf90f9cab53f95.jpg: 416x416 1 Caps, Done. (0.031s)
image 20/20 /content/yolov5/./test/images/cap_98_jpg.rf.28d796ab0d219b3ab51ca37cf6cc2945.jpg: 416x416 1 Caps, Done. (0.031s)
Results saved to runs/detect/exp
Done. (0.843s)
```


```
[ ]: import glob
from IPython.display import Image, display


for imageName in glob.glob('/content/yolov5/runs/detect/exp/*.jpg'): #assuming
    ↪JPG
    display(Image(filename=imageName))
    print("\n")
```























































Figure 17: Detected objects with label and probability